

PDLP: A Practical First-Order Method for Large-Scale Linear Programming



PDHG (Primal-Dual hybrid gradient)

Consider the following primal-dual LP:

$$\min c^T x \quad \max q^T y + l^T \lambda^+ - u^T \lambda^- \quad (1)$$

$$\text{s.t. } Ax = b \quad \text{s.t. } c - K^T y = \lambda \quad (2)$$

$$Gx \geq h \quad y_{1:m_1} \geq 0 \quad (3)$$

$$l \leq x \leq u \quad \lambda \in \Lambda \quad (4)$$

where $G \in \mathbb{R}^{m_1 \times n}$, $A \in \mathbb{R}^{m_2 \times n}$, $c \in \mathbb{R}^n$, $h \in \mathbb{R}^{m_1}$, $b \in \mathbb{R}^{m_2}$,
 $l \in (\mathbb{R} \cup \{-\infty\})^n$, $u \in (\mathbb{R} \cup \{+\infty\})^n$, $K^T = (G^T, A^T)$, $q^T = (h^T, b^T)$,
and

$$\Lambda = \Lambda_1 \times \cdots \times \Lambda_n, \quad \Lambda_i = \begin{cases} \{0\} & l_i = -\infty, u_i = +\infty \\ \mathbb{R}^- & l_i = -\infty, u_i \in \mathbb{R} \\ \mathbb{R}^+ & l_i \in \mathbb{R}, u_i = +\infty \\ \mathbb{R} & \text{otherwise} \end{cases}$$

Saddle-point problem

To solve it, we can consider the equivalent saddle-point problem:

$$\min_{x \in X} \max_{y \in Y} \mathcal{L}(x, y) = c^\top x - y^\top Kx + q^\top y \quad (2)$$

with $X = \{x \in \mathbb{R}^n : l \leq x \leq u\}$ and $Y = \mathbb{R}^{m_2} \times \mathbb{R}_+^{m_1}$.

PDHG

Constrained PDHG

LP problems usually have constraints, so we need constrained PDHG.

Primal update

$$x^{t+1} = \text{proj}_X(x^t - \tau(c - K^\top y^t))$$

Dual update

$$y^{t+1} = \text{proj}_Y(y^t + \sigma(q - K(2x^{t+1} - x^t))) .$$

τ and σ are step sizes for primal and dual. proj_Y and proj_X are projections for primal and dual, $\text{proj}_X(x) = \operatorname{argmin}_{z \in X} \|x - z\|_2$. Dual variable λ is not explicitly generated during PDHG, instead we compute

$$\lambda = \text{proj}_\Lambda(c - K^\top y)$$

The termination for PDHG relies on duality gap

Duality gap

Duality gap is the difference between primal and dual solutions. It can be measured with the Lagrangian:

$$\max_{\hat{z} \in Z} \{\mathcal{L}(x, \hat{y}) - \mathcal{L}(\hat{x}, y)\}.$$

For minimization problem, the optimal duality gap is always greater or equal 0, the it equals to 0 iff the strong duality holds (which obviously holds for LP).

The saddle point problem is to find a pair of solutions $z^* = (x^*, y^*)$ that $\mathcal{L}(x^*, y) \leq \mathcal{L}(x^*, y^*) \leq \mathcal{L}(x, y^*)$. Thus for saddle point problem constructed by LP, the theoretical optimal solution will be reached when the duality gap is 0.

Algorithm 1 PDHG-C

- 1: **Initialization**
 - 2: Pick y^0 and a feasible $x^0 \in X$, set $k \leftarrow 0$;
 - 3: **Choose step sizes**
 - 4: Choose step sizes τ_k and θ_k ;
 - 5: **repeat**
 - 6: **Updating**
 - 7: $x^{k+1} = \text{proj}_X(x^k - \tau(c - K^\top y^k));$
 - 8: $y^{k+1} = \text{proj}_Y(y^k + \sigma(q - K(2x^{k+1} - x^k)));$
 - 9: **Termination check**
 - 10: $k \leftarrow k + 1$;
 - 11: **until** a stopping criterion is satisfied;
 - 12: **Output:** x^k, y^k .
-

PDHG norm and step reparameterization

The step sizes τ and σ can be reparameterization as

$$\tau = \eta/\omega, \sigma = \eta\omega,$$

with $\eta \in (0, \infty)$ and $\omega \in (0, \infty)$. We call ω the primal weight, and η the step size now.

The PDHG norm then can be defined with primal weight ω :

$$\|z\|_\omega := \sqrt{\omega\|x\|_2^2 + \frac{\|y\|_2^2}{\omega}},$$

where z is the primal-dual vector (x, y) . The PDHG norm plays a role in theory for the PDHG and practice for the PDLP.

Algorithm 2 PDLP (after preconditioning and presolve)

- 1: **Input:** An initial solution $z^{0,0}$.
- 2: Initialize outer loop counter $n \leftarrow 0$, total iterations $k \leftarrow 0$, step size $\hat{\eta}^{0,0} \leftarrow 1/\|K\|_\infty$, primal weight $\omega^0 \leftarrow \text{InitializePrimalWeight}(c, q)$;
- 3: **repeat**
- 4: $t \leftarrow 0$;
- 5: **repeat**
- 6: $z^{n,t+1}, \eta^{n,t+1}, \hat{\eta}^{n,t+1} \leftarrow \text{AdaptiveStepOfPDHG}(z^{n,t}, \omega^n, \hat{\eta}^{n,t}, k)$;
- 7: $\bar{z}^{n,t+1} \leftarrow \frac{1}{\sum_{i=1}^{t+1} \eta^{n,i}} \sum_{i=1}^{t+1} \eta^{n,i} z^{n,i}$;
- 8: $z_c^{n,t+1} \leftarrow \text{GetRestartCandidate}(z^{n,t+1}, z_m^{n,t+1}, z^{n,0})$;
- 9: $t \leftarrow t + 1, k \leftarrow k + 1$;
- 10: **until** restart or termination criteria holds;
- 11: restart the outer loop $z^{n+1,0} \leftarrow z_c^{n,t}, n \leftarrow n + 1$;
- 12: $\omega^n \leftarrow \text{PrimalWeightUpdate}(z^{n,0}, z^{n-1,0}, \omega^{n-1})$;
- 13: **until** termination criteria holds;
- 14: **Output:** $z^{n,0}$.

Improvements

PDLP has some improvements to achieve better practical performance compare with PDHG.

- Preconditioning
- Adaptive restarts
- Adaptive step size
- Primal weight update
- Infeasibility detection

PDLP-Preconditioning

Preconditioning is a popular heuristic in optimization for improving the convergence of FOM (First-order method) by shrinking the condition number of coefficient matrix K .

Condition number

Condition number measures how much the output value of a function can change for a small change in the input argument (sensitivity to input's errors). For a linear equation system $Ax = b$, its condition number is $k(A) = \|A\| \cdot \|A^\dagger\|$, the norm often takes l_2 , the A^\dagger is the Moore-Penrose pseudoinverse of matrix A .

PDLP-Preconditioning

Many preconditioning methods contain matrix factorizations. Thus PDLP only consider using diagonal preconditioners: Pock-Chambolle and Ruiz. Preconditioning create a new LP instance. With diagonal matrices D_1 and D_2 , K is rescaled to $D_1KD_2 = \tilde{K} = (\tilde{G}, \tilde{A})$,
 $\tilde{x} = D_2^{-1}x$, $\tilde{c} = D_2c$, $(\tilde{b}, \tilde{h}) = D_1(b, h)$, $\tilde{u} = D_2^{-1}u$, $\tilde{l} = D_2^{-1}l$.

Pock-Chambolle

The preconditioners are defined by $(D_1)_{jj} = 1/\sqrt{\|K_{j,\cdot}\|_{2-\alpha}}$ for $j = 1, \dots, m_1 + m_2$ and $(D_2)_{ii} = 1/\sqrt{\|K_{\cdot,i}\|_\alpha}$ for $i = 1, \dots, n$. In PDLP we choose $\alpha = 1$.

Ruiz

The preconditioners are defined by $(D_1)_{jj} = 1/\sqrt{\|K_{j,\cdot}\|_\infty}$ for $j = 1, \dots, m_1 + m_2$ and $(D_2)_{ii} = 1/\sqrt{\|K_{\cdot,i}\|_\infty}$ for $i = 1, \dots, n$.

In particular, PDLP applies 10 iterations of Ruiz scaling and then Pock-Chambolle.



PDLP-Preconditioning

Consider the ill-conditioned 2×4 matrix:

$$A = \begin{bmatrix} 1 & 1 & 1 & 1 \\ 0 & \epsilon & 0 & \epsilon \end{bmatrix}, \quad \epsilon = 10^{-6}.$$

- **Condition Number:** Using SVD, singular values are $\sigma_1 \approx 2$, $\sigma_2 \approx \epsilon$.
Thus:

$$k(A) = \frac{\sigma_1}{\sigma_2} \approx \frac{2}{\epsilon} = 2 \times 10^6.$$

PDLP-Preconditioning

- **Pock-Chambolle Method ($\alpha = 1$):** Preconditioners use ℓ_1 -norm:

$$(D_1)_{jj} = 1/\sqrt{\|A_{j,\cdot}\|_1}, \quad (D_2)_{ii} = 1/\sqrt{\|A_{\cdot,i}\|_1}.$$

- $D_1 = \text{diag}(1, 0.9999995, 1, 0.9999995) \approx \text{diag}(1, 1, 1, 1)$,
 $D_2 = \text{diag}(0.5, 707.1068)$.
- Processed matrix: $D_2 A D_1 = \begin{bmatrix} 0.5 & 0.5 & 0.5 & 0.5 \\ 0 & 0.7071068 & 0 & 0.7071068 \end{bmatrix}$.
- Condition number: $\sigma_1 \approx 1.3066$, $\sigma_2 \approx 0.5412$, $k \approx 2.414$ (vs. 2×10^6).
- **Ruiz Method:** Preconditioners use ℓ_∞ -norm, iterative:

$$(D_1)_{jj} = 1/\sqrt{\|A_{j,\cdot}\|_\infty}, \quad (D_2)_{ii} = 1/\sqrt{\|A_{\cdot,i}\|_\infty}.$$

- First iteration: $D_1 = \text{diag}(1, 1, 1, 1)$, $D_2 = \text{diag}(1, 1000)$.
- Processed matrix: $D_2 A D_1 = \begin{bmatrix} 1 & 1 & 1 & 1 \\ 0 & 0.001 & 0 & 0.001 \end{bmatrix}$.
- Condition number: $\sigma_1 \approx 2$, $\sigma_2 \approx 0.0014142$, $k \approx 1414.21$ (vs. 2×10^6).

PDLP-Adaptive step size

Convergence promise

$$\eta \leq \frac{\|z^{k+1} - z^k\|_{\omega}^2}{2(y^{k+1} - y^k)^T K(x^{k+1} - x^k)}$$

Step size choice

- Classical choice:

$$\eta = \frac{1}{\|K\|_2}$$

PDLP will always converge if $\eta \leq \frac{1}{\|K\|_2}$. But this choice is "pessimistic" and requires the estimation of $\|K\|_2$.

- Adaptive choice: ALGO 3

PDLP-Adaptive step size

Algorithm 3 One step of PDHG using our step size heuristic

```
1: function ADAPTIVESTEPOFPDHG( $z^{n,t}, w^n, \eta^{n,t}, k$ )
2:    $(x, y) \leftarrow z^{n,t}, \eta \leftarrow \eta^{n,t};$ 
3:   for  $i \leftarrow 1, \dots, \infty$  do
4:      $x' \leftarrow \text{proj}_X \left( x - \frac{\eta}{w^n} (c - K^T y) \right);$ 
5:      $y' \leftarrow \text{proj}_Y \left( y + \eta w^n (q - K(2x' - x)) \right);$ 
6:      $\bar{\eta} \leftarrow \frac{\|(x-x', y-y')\|_{\omega^n}^2}{2(y'-y)^T K(x'-x)};$ 
7:      $\eta' \leftarrow \min \left( (1 - (k+1)^{-0.3}) \bar{\eta}, (1 + (k+1)^{-0.6}) \eta \right);$ 
8:     if  $\eta \leq \bar{\eta}$  then
9:       return  $(x', y'), \eta, \eta';$ 
10:    end if
11:     $\eta \leftarrow \eta';$ 
12:   end for
13: end function
```

PDLP-Adaptive step size

Note

- η will be reduced again if $\eta > \bar{\eta}$. The ALGO will be repeated then.
- $\bar{\eta} \geq \frac{1}{\|K\|_2}$ always holds, which make $\eta \geq \frac{1-o(1)}{\|K\|_2}$ holds as $k \rightarrow \infty$.
- The exponents -0.3 and -0.6 are heuristics, likely found experimentally, ensuring decrease is more aggressive than increase, and adjustments diminish as k grows, leading to stabilization.

PDLP-Primal weight update

$$\text{InitializePrimalWeight}(c, q) := \begin{cases} \frac{\|c\|_2}{\|q\|_2} & \|c\|_2, \|q\|_2 > \epsilon_{\text{zero}} \\ 1 & \text{otherwise} \end{cases},$$

where ϵ_{zero} is a small nonzero tolerance.

Algorithm 4 Primal weight update

```
1: function PRIMALWEIGHTUPDATE( $z^{n,0}, z^{n-1,0}, \omega^{n-1}$ )
2:    $\Delta_x^n \leftarrow \|x^{n,0} - x^{n-1,0}\|_2$ ,  $\Delta_y^n \leftarrow \|y^{n,0} - y^{n-1,0}\|_2$ ;
3:   if  $\Delta_x^n > \epsilon_{\text{zero}}$  and  $\Delta_y^n > \epsilon_{\text{zero}}$  then
4:     return  $\exp\left(\theta \log\left(\frac{\Delta_y^n}{\Delta_x^n}\right) + (1 - \theta) \log(\omega^{n-1})\right)$ ;
5:   else
6:     return  $\omega^{n-1}$ ;
7:   end if
8: end function
```



PDLP-Primal weight update

Insight of the update function

Heuristically we want to adjust the primal weight ω^n such that distance optimality in the primal and dual is almost the same, i.e.

$$\|(x^{n,t} - x^*, 0)\|_\omega \approx \|(0, y^{n,t} - y^*)\|_\omega,$$

where y^* and x^* are optimal solution for dual and primal. By definition of PDHG norm, we have

$$\|(x^{n,t} - x^*, 0)\|_\omega = \sqrt{\omega} \|x^{n,t} - x^*\|_2, \quad \|(0, y^{n,t} - y^*)\|_\omega = \frac{1}{\sqrt{\omega}} \|y^{n,t} - y^*\|_2.$$

To make the two terms equal yields $\omega = \|y^{n,t} - y^*\|_2 / \|x^{n,t} - x^*\|_2$. The function

$$\exp \left(\theta \log \left(\frac{\Delta_y^n}{\Delta_x^n} \right) + (1 - \theta) \log(\omega^{n-1}) \right)$$

estimate the ω term with Δ_x^n and Δ_y^n , and dampen the variation with log-scale and exponential smooth ($\theta \in [0, 1]$, PDLP takes $\theta = 0.5$).

PDLP-Infeasibility detection

Divergence behavior

When the LP is infeasible, PDLP iterates $z^k = (x^k, y^k)$ don't converge. Instead:

- They diverge towards infinity, but *along a specific direction*.
- This movement follows a **ray**: starting from some point z^* and extending infinitely in a direction $v = (v_x, v_y)$.
- The direction v is called the **infimal displacement vector**.

The infimal displacement vector

The key insight from Lu et al. (2021) connects the direction vector $v = (v_x, v_y)$ to infeasibility:

- If Primal LP is Infeasible \iff Dual component $v_y \neq 0$.
- If Dual LP is Infeasible \iff Primal component $v_x \neq 0$.
- If Both Feasible $\implies v = (0, 0)$.

PDLP-Infeasibility detection

Detection in PDLP

PDLP periodically estimates the direction v from the iterates $z^{n,t}$ by three sequences converging to v starting from the most recent restart:

- ① Difference of Iterates: $d^{n,t} = z^{n,t+1} - z^{n,t}$
- ② Normalized Iterates: $(z^{n,t} - z^{n,0})/t$
- ③ Normalized Average: $2(\bar{z}^{n,t} - \bar{z}^{n,0})/(t + 1)$

Note

Verifying all three sequences is often faster than relying on just one, since their convergence speed varies with the problem's geometry.

PDLP-Adaptive restart

Normalized duality gap

Recall that duality gap was defined as

$$\max_{\hat{z} \in Z} \{\mathcal{L}(x, \hat{y}) - \mathcal{L}(\hat{x}, y)\}.$$

However this function can lead to infinity if the problem is unbounded. So Lu et al. (2022) introduces normalized duality gap $\rho_r(z)$:

$$\rho_r(z) := \frac{\max_{\hat{z} \in W_r(z)} \{\mathcal{L}(x, \hat{y}), \mathcal{L}(\hat{x}, y)\}}{r},$$

where $W_r(z) := \{\hat{z} \in Z \mid \|z - \hat{z}\| \leq r\}$ is the ball centered at z with radius $r \in (0, \infty)$ intersected with the set Z , and $\|\cdot\|$ is a carefully selected semi-norm on Z . In PDLP, the authors redefined $\mu_n(z, z_{ref})$ as the normalized duality gap at z with radius $\|z - z_{ref}\|_{\omega^n}$, i.e.

$$\mu_n(z, z_{ref}) = \rho_{\|z - z_{ref}\|_{\omega^n}}^n(z).$$

PDLP-Adaptive restart

Restart Candidate

GetRestartCandidate($z^{n,t+1}, \bar{z}^{n,t+1}, z^{n,0}$) :=
$$\begin{cases} z^{n,t+1} & \mu_n(z^{n,t+1}, z^{n,0}) < \mu_n(\bar{z}^{n,t+1}, z^{n,0}) \\ \bar{z}^{n,t+1} & \text{otherwise} \end{cases}.$$

Obviously this is a heuristic for greedy step size choices.

PDLP-Adaptive restart

Restart criteria

① Sufficient decay in normalized duality gap:

$$\mu_n(z_c^{n,t+1}, z^{n,0}) \leq \beta_{\text{sufficient}} \mu_n(z^{n,0}, z^{n-1,0})$$

② Necessary decay + no local progress in normalized duality gap:

$$\mu_n(z_c^{n,t+1}, z^{n,0}) \leq \beta_{\text{necessary}} \mu_n(z^{n,0}, z^{n-1,0})$$

$$\mu_n(z_c^{n,t+1}, z^{n,0}) > \mu_n(z_c^{n,t}, z^{n,0})$$

③ Long inner loop: $t \geq \beta_{\text{artificial}} k$.

The parameters are $\beta_{\text{necessary}} \in (0, 1)$, $\beta_{\text{sufficient}} \in (0, \beta_{\text{necessary}})$ and $\beta_{\text{artificial}} \in (0, 1)$. PDLP uses $\beta_{\text{necessary}} = 0.9$, $\beta_{\text{sufficient}} = 0.1$ and $\beta_{\text{artificial}} = 0.5$.

PDLP-Adaptive restart

Let's break down the insight of the three criteria.

- ➊ **Sufficient decay in normalized duality gap:** The duality gap for n iterate is considered sufficiently decayed when the progress of the n iterate is detected to have reached the $\beta_{\text{sufficient}}$ progress of the $n - 1$ iterate. Restart at this point.
- ➋ **Necessary decay + no local progress in normalized duality gap:** The duality gap for n iterate is considered necessarily decayed when the progress of the n iterate is detected to have reached the $\beta_{\text{necessary}}$ progress of the $n - 1$ iterate. Restart will happen if the duality gap oscillates at this point.
- ➌ **Long inner loop:** Restart when internal loop time is too long.

PDLP-Adaptive restart

Sharpness

For a given function f , we say f is α -sharp if it holds for any $z \in Z$ that

$$f(z) - f^* \geq \alpha \mathbf{dist}(z, Z^*)^p,$$

where Z^* is the set of optima of the function f , f^* is the optimal function value, and $\mathbf{dist}(z, Z) := \inf_{\hat{z} \in Z} \|z - \hat{z}\|$. $p = 1$ means weak sharpness, $p = 2$ means strong sharpness.

Sharpness is related to the speed of convergence near the optimal solution of function f .

Theoretical guarantee

Lu et al. (2022) gives a restart scheme for the primal-dual class of algorithms, and also gives a convergence analysis of the algorithms based on sharpness.

- LP-constructed primal-dual problems are shown to be α -sharp within $W_R(0)$, where $W_R(0) = \{\hat{z} \in Z \mid |\hat{z}| \leq R\} \cap Z$.
- For restarted PDHG, there exists a constant $R > 0$ such that $z^{n,0} \in W_R(z^{0,0})$ for all $n \in \mathbb{N}$.
- If the primal-dual problem α -sharp on the set S containing all $z^{n,0}$, the distance to primal-dual solution decays linearly.

PDLP-Termination criteria

- **Duality gap:**

$$|q^T y + I^T \lambda^+ - u^T \lambda^- - c^T x| \leq \epsilon(1 + |q^T y + I^T \lambda^+ - u^T \lambda^-| + |c^T x|)$$

- **Primal feasibility:**

$$\left\| \begin{pmatrix} Ax - b \\ (h - Gx)^+ \end{pmatrix} \right\|_2 \leq \epsilon(1 + \|q\|_2)$$

- **Dual feasibility:**

$$\|c - K^T y - \lambda\|_2 \leq \epsilon(1 + \|c\|_2)$$

where $\epsilon \in (0, \infty)$ is the termination tolerance. PDLP sets $\epsilon = 10^{-4}$ for moderately accurate solutions and $\epsilon = 10^{-8}$ as a benchmark for high-quality solutions.



Algorithm 5 PDLP (after preconditioning and presolve)

```

1: Input: An initial solution  $z^{0,0}$ .
2: Initialize outer loop counter  $n \leftarrow 0$ , total iterations  $k \leftarrow 0$ , step size
    $\hat{\eta}^{0,0} \leftarrow 1/\|K\|_\infty$ , primal weight  $\omega^0 \leftarrow \text{InitializePrimalWeight}(c, q)$ ;
3: repeat
4:    $t \leftarrow 0$ ;
5:   repeat
6:      $z^{n,t+1}, \eta^{n,t+1}, \hat{\eta}^{n,t+1} \leftarrow \text{AdaptiveStepOfPDHG}(z^{n,t}, \omega^n, \hat{\eta}^{n,t}, k)$ ;
7:      $\bar{z}^{n,t+1} \leftarrow \frac{1}{\sum_{i=1}^{t+1} \eta^{n,i}} \sum_{i=1}^{t+1} \eta^{n,i} z^{n,i}$ ;
8:      $z_c^{n,t+1} \leftarrow \text{GetRestartCandidate}(z^{n,t+1}, z_m^{n,t+1}, z^{n,0})$ ;
9:      $t \leftarrow t + 1, k \leftarrow k + 1$ ;
10:    until restart or termination criteria holds;
11:    restart the outer loop  $z^{n+1,0} \leftarrow z_c^{n,t}, n \leftarrow n + 1$ ;
12:     $\omega^n \leftarrow \text{PrimalWeightUpdate}(z^{n,0}, z^{n-1,0}, \omega^{n-1})$ ;
13: until termination criteria holds;
14: Output:  $z^{n,0}$ .

```

Adaptive restart based on KKT error

cuPDLP uses the KKT error defined below as a proxy to the normalized duality gap for restarting since KKT error is more computational GPU-friendly.

$$KKT_{\omega}(z) = \sqrt{d^2 + \omega^2 \left\| \begin{pmatrix} Ax - b \\ (h - Gx)^+ \end{pmatrix} \right\|_2^2 + \frac{1}{\omega^2} \|c - K^T y - \lambda\|_2^2}$$

$$d = q^T y + I^T \lambda^+ - u^T \lambda^- - c^T x$$

KKT error restart criteria

- ① Sufficient decay in KKT error:

$$\text{KKT}_{\omega^n}(z_c^{n,t+1}) \leq \beta_{\text{sufficient}} \text{KKT}_{\omega^n}(z^{n,0})$$

- ② Necessary decay + no local progress in KKT error:

$$\text{KKT}_{\omega^n}(z_c^{n,t+1}) \leq \beta_{\text{necessary}} \text{KKT}_{\omega^n}(z^{n,0})$$

$$\text{KKT}_{\omega^n}(z_c^{n,t+1}) > \text{KKT}_{\omega^n}(z_c^{n,t})$$

- ③ Long inner loop: $t \geq \beta_{\text{artificial}} k$.

Performance on selected large-scale problems

Source	Instance	<i>m</i>	<i>n</i>	<i>nnz</i>	COPT	cuPDLP-C
Koch et al. [8]	zib03	19,731,970	29,128,799	104,422,573	16.5 (h)	916
Pagerank	rand_1m_nodes	1,000,001	1,000,000	7,999,982	-	3.56
	rand_10m_nodes	10,000,001	10,000,000	79,999,982	-	44.22
	com-livejournal	3,997,963	3,997,962	77,358,302	-	21.07
	soc-livejournal1	4,847,572	4,847,571	78,170,533	-	22.26
Unit Com.	ds1	641,037	659,145	21,577,566	592	81
	ds2	641,037	659,145	21,577,566	606	108
Supply-chain	inv-10	4,035,449	3,758,458	15,264,380	-	1636
	inv-20	8,368,795	7,810,584	31,718,673	-	1157
	inv-40	13,186,756	12,066,105	49,528,729	-	6032
	inv-60	16,227,780	14,544,689	60,372,404	-	11102
QAP [1]	wil50				-	96
	lipa50a				-	71
	lipa50b	3,437,600	6,252,500	19,125,000	-	66
	tai50a				-	64
	tai50b				-	437

Figure: Performance on selected large-scale problems

Partial results of Mittelmann feasible point instances

Method	COPT	cuPDLP-C (COPT)		cuPDLP-C (CLP)		cuPDLP-C	
	10^{-8}	10^{-4}	10^{-8}	10^{-4}	10^{-8}	10^{-4}	10^{-8}
L1.sixm250obs	0.79	6.87	7.99	3.27	5.29	5.34	6.59
Linf_520c	3.07	1253.58	3600.00	11.07	3600.00	13.29	3600.00
a2864-99blp	0.03	0.26	0.34	0.35	0.52	0.53	0.92
bdry2	6.21	3600.00	3600.00	3600.00	3600.00	3600.00	3600.00
cont1	0.09	238.41	602.34	289.77	1337.23	258.78	621.90
cont11	0.19	386.36	3419.51	227.04	3600.00	497.35	2352.79
datt256	0.79	0.17	0.32	0.12	0.24	0.12	0.28
dlr1	26.58	13.36	3600.00	408.29	3600.00	3600.00	3600.00
ex10	0.61	0.06	0.09	0.06	0.10	0.07	0.15
fhnw-binschedule1	21.60	17.50	20.68	55.64	114.75	70.22	75.67
fome13	0.56	1.41	22.23	3.74	11.67	1.54	7.23
graph40-40-1rand	0.02	0.02	0.02	0.05	0.12	0.04	0.09
irish-electricity	11.41	11.63	263.26	5.79	250.40	11.31	3600.00
neos	6.70	22.93	245.26	21.13	212.24	32.58	331.62
neos3	0.11	0.26	0.49	0.39	0.82	0.39	0.80
neos-3025225-shelon	8.51	6.35	6.63	2.79	2.94	2.83	2.97
neos-5052403-cygnets	3.64	9.13	22.56	2.29	9.17	2.21	10.43
neos-5251015-ogosta	1.75	0.73	1.06	0.49	1.00	0.25	2.27
ns1687037	3.02	3600.00	3600.00	237.89	3600.00	3600.00	3600.00
ns1688926	6.55	3600.00	3600.00	3600.00	3600.00	3600.00	3600.00

Figure: Partial results of Mittelmann feasible point instances

Overall performance (SGM10) on different datasets

Dataset	Optimizer	Presolver	Tol.	SGM10	Solved
MIPLIB (383)	COPT	-	10^{-8}	3.11	383
			10^{-4}	5.43	379
	cuPDLP-C	COPT	10^{-4}	18.53	369
			10^{-8}	7.95	372
		CLP	10^{-4}	21.89	362
			10^{-8}	10.28	370
	cuPDLP.jl	No Presolve	10^{-4}	27.15	359
			10^{-8}	17.49	370
		No Presolve	10^{-4}	35.69	355
			10^{-8}	48	48
Mittelmann (49)	COPT	-	10^{-8}	25.29	46
			10^{-4}	110.22	41
	cuPDLP-C	COPT	10^{-4}	33.97	45
			10^{-8}	125.95	38
		CLP	10^{-4}	57.54	43
			10^{-8}	172.98	39

Figure: Overall performance (SGM10) on different datasets

Shifted geometric mean: $SGM := (\prod_{i=1}^n (t_i + \Delta))^{1/n} - \Delta$, where t_i is the solving time for the i th instance and $\Delta = 10$.